


```
In [55]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.dummy import DummyClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split
```

```
In [56]: data = pd.read_csv('../datasets/high_diamond_ranked_10min.csv')
```

```
In [57]: plt.figure(figsize=(20,10))
sns.heatmap(data.corr(), cmap='RdBu_r', annot=True)
```

```
Out[57]: <matplotlib.axes._subplots.AxesSubplot at 0x14b94b1ee10>
```



```
In [ ]:
```

```
In [58]: data = data.drop(columns=['gameId', 'redGoldPerMin', 'redKills', 'redDeaths',
'blueGoldPerMin',
'blueCSPerMin', 'redCSPerMin', 'redFirstBlood', 'redGoldDiff',
'redExperienceDiff', 'blueTotalGold', 'redTotalGold',
'blueTotalExperience', 'redTotalExperience'])

data['blueWardsPlacedDiff'] = data['blueWardsPlaced'] - data['redWardsPlaced']
data['blueWardsDestroyedDiff'] = data['blueWardsDestroyed'] - data['redWardsDestroyed']
data['blueAvgLevelDiff'] = data['blueAvgLevel'] - data['redAvgLevel']
data['blueAssistsDiff'] = data['blueAssists'] - data['redAssists']
data['blueTotalMinionsKilledDiff'] = data['blueTotalMinionsKilled'] - data['redTotalMinionsKilled']
data['blueTotalJungleMinionsKilledDiff'] = data['blueTotalJungleMinionsKilled'] - data['redTotalJungleMinionsKilled']
data['blueEliteMonstersDiff'] = data['blueEliteMonsters'] - data['redEliteMonsters']
data['blueDragonsDiff'] = data['blueDragons'] - data['redDragons']
data['blueHeraldsDiff'] = data['blueHeralds'] - data['redHeralds']
data['blueTowersDestroyedDiff'] = data['blueTowersDestroyed'] - data['redTowersDestroyed']
```

```
In [59]: data = data.drop(columns= ['blueWardsPlaced', 'redWardsPlaced', 'blueWardsDestroyed', 'redWardsDestroyed',
'blueAvgLevel', 'redAvgLevel', 'blueAssists', 'redAssists', 'blueTotalMinionsKilled',
'redTotalMinionsKilled', 'blueTotalJungleMinionsKilled', 'redTotalJungleMinionsKilled',
'blueEliteMonsters', 'redEliteMonsters', 'redDragons', 'blueDragons', 'blueHeralds',
'redHeralds', 'blueTowersDestroyed', 'redTowersDestroyed'])
data.head()
```

Out[59]:

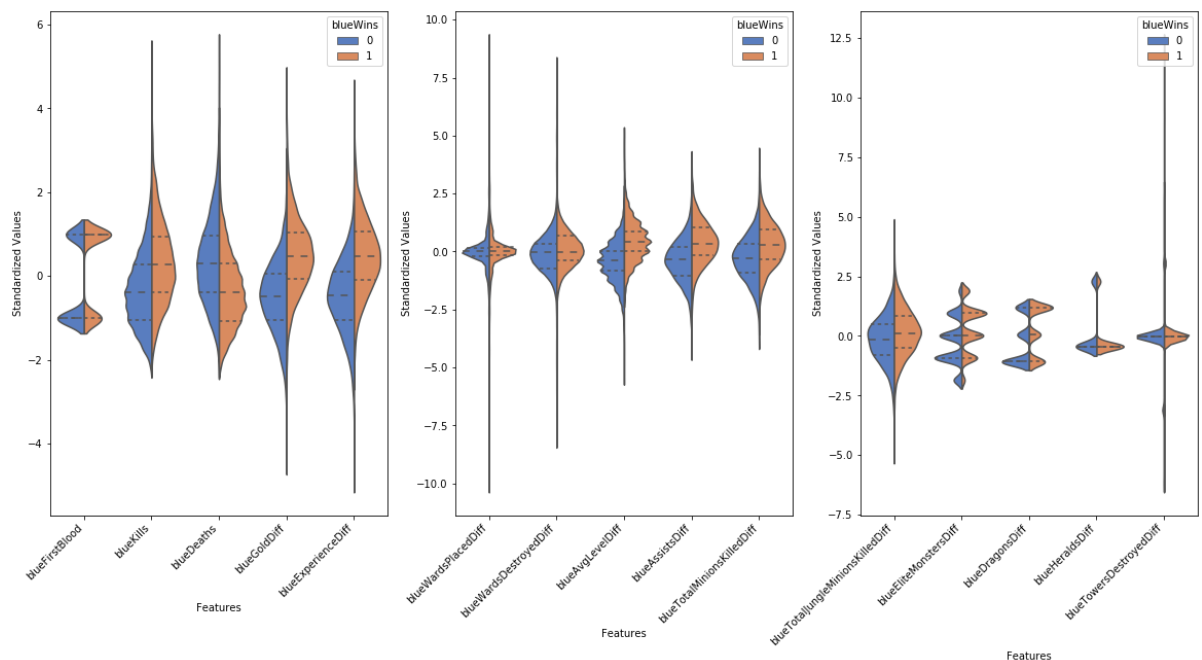
	blueWins	blueFirstBlood	blueKills	blueDeaths	blueGoldDiff	blueExperienceDiff	blueWardsPlacedDiff
0	0	1	9	6	643	-8	
1	0	0	5	5	-2908	-1173	
2	0	0	7	11	-1172	-1033	
3	0	0	4	5	-1321	-7	
4	0	0	6	6	-1004	230	

```

In [60]: def plot_violinplot(df, ax_key):
    df = pd.melt(df, id_vars='blueWins', var_name='Features', value_name='Standardized Values')
    sns.violinplot(x='Features', y='Standardized Values', hue='blueWins', data=df, split=True,
                   inner='quart', ax=ax[ax_key], palette='muted')
    fig.autofmt_xdate(rotation=45)

fig, ax = plt.subplots(1,3,figsize=(20,10))
df = data.loc[:, data.columns != 'blueWins']
df_std = StandardScaler().fit_transform(df)
df_std = pd.DataFrame(data = df_std, columns = df.columns)
df = pd.concat([data.blueWins, df_std.iloc[:, 0:5]], axis=1)
plot_violinplot(df,0)
df2 = pd.concat([data.blueWins, df_std.iloc[:, 5:10]], axis=1)
plot_violinplot(df2,1)
df3 = pd.concat([data.blueWins, df_std.iloc[:, 10:]], axis=1)
plot_violinplot(df3,2)
plt.show()

```




```
In [61]: X = data.loc[:, data.columns != 'blueWins']
y = data['blueWins']
X_train2, X_test2, y_train2, y_test2 = train_test_split(X, y, test_size=0.3, random_state=42)
X_train, X_test, y_train, y_test = train_test_split(X_train2, y_train2, test_size=0.3, random_state=43)
X_train
```

Out[61]:

	blueFirstBlood	blueKills	blueDeaths	blueGoldDiff	blueExperienceDiff	blueWardsPlacedDiff
3189	1	6	7	457	-383	0
5249	1	10	5	1680	1254	-1
785	0	3	3	-5	-269	-21
8903	1	10	8	-1125	-1326	3
1092	1	10	5	4534	1707	70
...
4300	0	10	8	496	187	0
7619	1	6	6	-695	-683	3
5857	0	7	6	1835	393	2
5409	1	12	7	2473	565	-4
3368	0	0	9	-4912	-4457	23

4840 rows × 15 columns



In []:

```
In [62]: from sklearn.metrics import classification_report
from matplotlib.colors import ListedColormap
```

```
In [63]: def plot_decision_regions(X, y, classifier, resolution=0.02):
# setup marker generator and color map
markers = ('s', 'x', 'o', '^', 'v')
colors = ('red', 'blue', 'lightgreen', 'gray', 'cyan')
cmap = ListedColormap(colors[:len(np.unique(y))])

# plot the decision surface
x1_min, x1_max = X[:, 0].min() - 1, X[:, 0].max() + 1
x2_min, x2_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution),
                        np.arange(x2_min, x2_max, resolution))
Z = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)
Z = Z.reshape(xx1.shape)
plt.contourf(xx1, xx2, Z, alpha=0.3, cmap=cmap)
plt.xlim(xx1.min(), xx1.max())
plt.ylim(xx2.min(), xx2.max())

# plot class examples
for idx, cl in enumerate(np.unique(y)):
    plt.scatter(x=X[y == cl, 0],
                y=X[y == cl, 1],
                alpha=0.8,
                c=colors[idx],
                marker=markers[idx],
                label=cl,
                edgecolor='black')

plt.xlabel('feature 1')
plt.ylabel('feature 2')
plt.legend()
```

```
In [64]: clfL = DSELinearClassifier(activation='Logistic')
```

```

In [65]: from sklearn.base import BaseEstimator
from sklearn.base import ClassifierMixin
from sklearn.preprocessing import LabelEncoder
from sklearn.base import clone
from sklearn.pipeline import _name_estimators
import numpy as np
import operator
class MajorityVoteClassifier(BaseEstimator,
                           ClassifierMixin):
    """ A majority vote ensemble classifier

    Parameters
    -----
    classifiers : array-like, shape = [n_classifiers]
        Different classifiers for the ensemble

    vote : str, {'classlabel', 'probability'}
        Default: 'classlabel'
        If 'classlabel' the prediction is based on
        the argmax of class labels. Else if
        'probability', the argmax of the sum of
        probabilities is used to predict the class label
        (recommended for calibrated classifiers).

    weights : array-like, shape = [n_classifiers]
        Optional, default: None
        If a list of `int` or `float` values are
        provided, the classifiers are weighted by
        importance; Uses uniform weights if `weights=None`.

    """
    def __init__(self, classifiers,
                 vote='classlabel', weights=None):

        self.classifiers = classifiers
        self.named_classifiers = {key: value for
                                   key, value in
                                   _name_estimators(classifiers)}

        self.vote = vote
        self.weights = weights

    def fit(self, X, y):
        """ Fit classifiers.

        Parameters
        -----
        X : {array-like, sparse matrix},
            shape = [n_examples, n_features]
            Matrix of training examples.

        y : array-like, shape = [n_examples]
            Vector of target class labels.

        Returns
        -----
        self : object

```

```

"""
if self.vote not in ('probability', 'classlabel'):
    raise ValueError("vote must be 'probability'"
                     "or 'classlabel'; got (vote=%r)"
                     % self.vote)
if self.weights and len(self.weights) != len(self.classifiers):
    raise ValueError("Number of classifiers and weights"
                     "must be equal; got %d weights,"
                     "%d classifiers"
                     % (len(self.weights),
                        len(self.classifiers)))
# Use LabelEncoder to ensure class labels start
# with 0, which is important for np.argmax
# call in self.predict
self.lablenc_ = LabelEncoder()
self.lablenc_.fit(y)
self.classes_ = self.lablenc_.classes_
self.classifiers_ = []
for clf in self.classifiers:
    fitted_clf = clone(clf).fit(X,
                               self.lablenc_.transform(y))
    self.classifiers_.append(fitted_clf)
return self

def predict(self, X):
    """ Predict class labels for X.

    Parameters
    -----
    X : {array-like, sparse matrix},
        Shape = [n_examples, n_features]
        Matrix of training examples.

    Returns
    -----
    maj_vote : array-like, shape = [n_examples]
        Predicted class labels.

    """
    if self.vote == 'probability':
        maj_vote = np.argmax(self.predict_proba(X), axis=1)
    else: # 'classlabel' vote

        # Collect results from clf.predict calls
        predictions = np.asarray([clf.predict(X)
                                for clf in
                                self.classifiers_]).T

        maj_vote = np.apply_along_axis(lambda x: np.argmax(
            np.bincount(x,
                        weights=self.weights)),
            axis=1,
            arr=predictions)
    maj_vote = self.lablenc_.inverse_transform(maj_vote)
    return maj_vote

```

```

def predict_proba(self, X):
    """ Predict class probabilities for X.

    Parameters
    -----
    X : {array-like, sparse matrix},
        shape = [n_examples, n_features]
        Training vectors, where
        n_examples is the number of examples and
        n_features is the number of features.

    Returns
    -----
    avg_proba : array-like,
        shape = [n_examples, n_classes]
        Weighted average probability for
        each class per example.

    """
    probas = np.asarray([clf.predict_proba(X)
                        for clf in self.classifiers_])
    avg_proba = np.average(probas, axis=0,
                          weights=self.weights)
    return avg_proba

def get_params(self, deep=True):
    """ Get classifier parameter names for GridSearch"""
    if not deep:
        return super(MajorityVoteClassifier,
                    self).get_params(deep=False)
    else:
        out = self.named_classifiers.copy()
        for name, step in self.named_classifiers.items():
            for key, value in step.get_params(
                deep=True).items():
                out['%s__%s' % (name, key)] = value
        return out

```

```

In [66]: clfl.fit(X_train, y_train, learning_rate=0.001)
y_pred = clfl.predict(X_test)
print(classification_report(y_test, y_pred))
#plot_decision_regions(X_test, y_test, clfl)

```

	precision	recall	f1-score	support
0	0.75	0.76	0.75	1073
1	0.74	0.73	0.74	1002
accuracy			0.75	2075
macro avg	0.75	0.75	0.75	2075
weighted avg	0.75	0.75	0.75	2075

C:\Users\XChen\Anaconda3\lib\site-packages\ipykernel_launcher.py:71: RuntimeWarning: divide by zero encountered in log


```
In [33]: from sklearn.pipeline import Pipeline

clf1 = LogisticRegression(random_state=42)
clf2 = DecisionTreeClassifier(random_state=42)
clf3 = KNeighborsClassifier()
clf4 = RandomForestClassifier(random_state=42)
clf5 = SVC(random_state=42, probability = True)

pipe1 = Pipeline([['sc', StandardScaler()], ['clf', clf1]])
pipe3 = Pipeline([['sc', StandardScaler()], ['clf', clf3]])
pipe5 = Pipeline([['sc', StandardScaler()], ['clf', clf5]])
mv_clf = MajorityVoteClassifier(classifiers=[pipe1, clf2, pipe3, clf4, pipe5])

clf_labels = ['Logistic regression', 'Decision tree', 'KNN', 'RandomForestClassifier', 'SVM', 'mv_clf']

print('10-fold cross validation:\n')
for clf, label in zip([pipe1, clf2, pipe3, clf4, pipe5, mv_clf], clf_labels):
    scores = cross_val_score(estimator=clf, X=X_train, y=y_train, cv=10, scoring='roc_auc')
    print("ROC AUC: %0.2f (+/- %0.2f) [%s]"
          % (scores.mean(), scores.std(), label))
```

10-fold cross validation:

```
ROC AUC: 0.80 (+/- 0.03) [Logistic regression]
ROC AUC: 0.62 (+/- 0.03) [Decision tree]
ROC AUC: 0.73 (+/- 0.03) [KNN]
ROC AUC: 0.78 (+/- 0.03) [RandomForestClassifier]
ROC AUC: 0.78 (+/- 0.03) [SVM]
ROC AUC: 0.77 (+/- 0.03) [mv_clf]
```

In []: